

# From Python to C++: A Journey to Efficient DBSCAN





Student: Wei-Lun Chiu, Jhao-Ting Chen  
Carnegie Mellon University, Pittsburgh, PA




## DBSCAN: What, Why, How?

- Density-Based Spatial Clustering of Applications with Noise
- Connectivity-based: it produce desired results



- Grid-based DBSCAN Visualization ( $\epsilon$  is minimum connectivity threshold)



## Poor Parallel Performance in Python



## Analysis of Python Parallelism

- The main obstacle: Global Interpreter Lock (GIL)
  - No multi-thread parallelism allowed!
- PyOMP
  - Generate machine code. Super-fast (near C-level performance).
  - Restrictions:
    - No dynamic array, set, etc. No complex data structure
    - No recursive. No incomplete type. (No tree allowed!)
- Multiprocessing Module
  - Generates processes. Each process has its own GIL.
  - Too much overhead, too slow. Only 1.7x speedup with 6 threads.

## Analyzing Serial Runtimes in C++

Procedure	Runtime(s)	Percentage
Gridify	0.011691	0.9538%
Mark in-grid cores	$1.68 \times 10^{-6}$	0.0001%
Mark out-grid cores	0.00318535	0.2599%
Expand	1.21085	98.7861%

Table 1: Runtime and percentage breakdown for each step of the procedure


## C++ Concurrency API for Parallelism

- Part of C++ Standard Library since C++11
- Parallel implementation of expand in grid-based DBSCAN using **lock-free** UNION-FIND data structure for neighboring cell connectivity.

## Workload Balancing with Work-Stealing

- Observed workload imbalance (**Red line** in the following figure)
- Implement a work-stealing mechanism to balance workload (**yellow line**)
  - Thread-Specific work queues and **fine-grained locks**
  - Non-busy threads attempt to steal work from  $(\text{treadID}+1) \% \text{ntreads}$ .

OpenMP, Concurrency, and Concurrency + stealing




## Performance Analysis with Perf

- Memory allocation/release (45%)
- Floating point arithmetic (28%)

## Floating Point Optimization with AVX

- Optimize floating point arithmetic with **AVX** instructions to achieve **12.79x** speedup.

OpenMP, Concurrency + stealing, and Concurrency + stealing + AVX



#Ts	OpenMP	Concurrency + steal	Con. + steal + AVX
1	60.5908 (0.99x)	59.9386 (1.00x)	33.683 (1.78x)
4	45.5341 (1.32x)	25.7059 (2.33x)	11.5379 (5.20x)
8	24.6547 (2.43x)	16.8995 (3.55x)	6.08641 (9.86x)
16	19.4427 (3.09x)	15.2388 (3.94x)	5.57829 (10.76x)
32	20.3203 (2.95x)	16.4425 (3.65x)	4.69244 (12.79x)
64	17.8841 (3.35x)	16.1772 (3.71x)	5.05127 (11.88x)
78	19.089 (3.14x)	15.9645 (3.76x)	4.80074 (12.50x)

Table 2: Execution times (in seconds) for various thread counts and parallelization strategies